

EDA, Missingness and Hypothesis Analysis of Congressman Stock Trades

Kevin Morales Nguyen - A17186624

Summary of Findings

Introduction

This dataset primarily contains information relating to stock trades made by members of the US congress. The majority of the data is categorical which had great impact on how to approach analysis. This analysis is primarily focused in inquiring further about how party affiliation is related to various other variables part of transactions.

Cleaning and EDA

It was observed that the owner variable had many values missing and the missingness mechanism was to be further explored. Cleaning involved converting the dates formatting to timestamps, making sure to replace user nan values with appropriate nan values for analysis. The data cleaning also involved manual cleanup of representative names in order to merge with a second dataset that contained party affiliation which was not part of the original dataset. EDA revealed that the two biggest stocks held amongst both republicans and democrats was Apple and Microsoft, further exploration showed that republicans tended to trade older technology and traditional energy companies compared to democrats who traded newer technology, solar energy and large holdings companies. Within the data set it was observed that democrats traded about 63% of the trades and republicans made up 36% of the trades, further more when analysing the amount traded by party it was found that democrats engaged in many more high usd value trades that were at or exceeded \$500,000.

Assessment of Missingness

It was found that the Owner variable may be NMAR, but MAR permutation testing showed that the missingness of owner is in fact MAR dependent on amount, type, and party variables.

Hypothesis Test

When hypothesizing about the proportion of trades based on party affiliation, a null hypothesis in which party affiliates traded at equal proportions was rejected in favor of a alternative hypothesis that democrats tend to initiate more trades than republicans.

Code

```
In [207]: ▶ import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd
import seaborn as sns
import requests
import bs4
```

Cleaning and EDA

Load the data and take an initial look

```
In [208]: ▶ raw_data = pd.read_csv('all_transactions.csv')
raw_data.head()
```

Out[208]:

	disclosure_year	disclosure_date	transaction_date	owner	ticker	asset_description	
0	2021	10/04/2021	2021-09-27	joint	BP	BP plc	pu
1	2021	10/04/2021	2021-09-13	joint	XOM	Exxon Mobil Corporation	pu
2	2021	10/04/2021	2021-09-10	joint	ILPT	Industrial Logistics Properties Trust - Common...	pu
3	2021	10/04/2021	2021-09-28	joint	PM	Phillip Morris International Inc	pu
4	2021	10/04/2021	2021-09-17	self	BLK	BlackRock Inc	sale

Okay merging political party on names might be a good idea, let's get clean up names column by removing 'Hon.' from all names, also by default if there was no value in csv is was loaded into dataframe with np.nan float but some values were loaded with '--' so let's change those to np.nan floats across all columns and rows and lets check to make sure all columns have appropriate data

```
In [299]: ▶ raw_data = pd.read_csv('all_transactions.csv')
#raw_data

raw_data['representative'] = raw_data['representative']
    .apply(lambda x: x.replace('Hon. ', ''))
#raw_data[:50]

raw_data_nan = raw_data.replace('--', np.nan)
#raw_data_nan['ticker'].value_counts()
```

The transaction_date has a different format to where we can't just easily convert it to timestamp so we let's make a helper function to prep is to that we can call to_datetime on cleaned series

```
In [210]: ▶ def convert_date(x):
           return x[5:7] + '/' + x[8:] + '/' + x[0:4]
```

Convert the transaction and disclose dates to timestamps

```
In [211]: ▶ raw_data_nan['disclosure_date'] = pd.to_datetime(
           raw_data_nan['disclosure_date'])

           raw_data_nan['transaction_date'] = raw_data_nan['transaction_date']
           .apply(convert_date)
           raw_data_nan['transaction_date'] = pd.to_datetime(
           raw_data_nan['transaction_date'])
           raw_data_nan
           #raw_data_nan['transaction_date'] = pd.to_datetime(raw_data_nan['tra
           #raw_data_nan
```

...
14269	2020	2020-06-10	2020-04-09	NaN	SWK	Stanley Black & Decker, Inc
14270	2020	2020-06-10	2020-04-09	NaN	USB	U.S. Bancor
14271	2020	2020-06-10	2020-03-13	NaN	BMJ	Bristol-Myer Squibb Compan
14272	2020	2020-06-10	2020-03-13	NaN	LLY	Eli Lilly an Compan
14273	2020	2020-06-10	2020-03-13	NaN	DIS	Walt Disne Compan

14274 rows × 12 columns

Let's go over a couple columns and check values to make sure different forms of nans aren't showing up

```
In [216]: ▶ print(raw_data_nan['disclosure_year'].value_counts())
           print('sum:', raw_data_nan['disclosure_year'].value_counts().sum())

           2020    7379
           2021    5520
           2022    1375
           Name: disclosure_year, dtype: int64
           sum: 14274
```

Overall data types, leaving disclosure year as int seems fine because it can be used as a broad

categorical variable, more refined data with months and days have been converted to datetimes, the rest are fine as strings in cludeing amount. The amount has a coarse granularity because it only captures ranges no a specific monetary amount so we will have to use it as a categorical variable. We could extract days and months for further accuracy if we want.

In [217]: `raw_data_nan.dtypes`

```
Out[217]: disclosure_year          int64
disclosure_date          datetime64[ns]
transaction_date         datetime64[ns]
owner                    object
ticker                   object
asset_description         object
type                     object
amount                   object
representative            object
district                 object
ptr_link                 object
cap_gains_over_200_usd    bool
dtype: object
```

Verify datetime data type

In [218]: `raw_data_nan['disclosure_date'][0]`

```
Out[218]: Timestamp('2021-10-04 00:00:00')
```

In [219]: `raw_data_nan`

...
14269	2020	2020-06-10	2020-04-09	NaN	SWK	Stanley Black & Decker, Inc
14270	2020	2020-06-10	2020-04-09	NaN	USB	U.S. Bancor
14271	2020	2020-06-10	2020-03-13	NaN	BMJ	Bristol-Myer Squibb Compan
14272	2020	2020-06-10	2020-03-13	NaN	LLY	Eli Lilly an Compan
14273	2020	2020-06-10	2020-03-13	NaN	DIS	Walt Disne Compan

Just going to check values to make sure nothing weird appears, or variation of nan value

```
In [220]: raw_data_nan['disclosure_year'].value_counts()
```

```
Out[220]: 2020    7379
          2021    5520
          2022    1375
          Name: disclosure_year, dtype: int64
```

```
In [221]: raw_data_nan['disclosure_date'].value_counts()
```

```
Out[221]: 2020-04-27    581
          2020-04-16    272
          2021-03-16    199
          2020-06-15    165
          2021-10-08    163
          ...
          2020-08-11      1
          2020-12-24      1
          2021-11-23      1
          2021-05-27      1
          2022-04-27      1
          Name: disclosure_date, Length: 606, dtype: int64
```

```
In [222]: raw_data_nan['transaction_date'].value_counts()
```

```
Out[222]: 2020-03-18    292
          2019-06-24    204
          2021-02-16    194
          2020-11-13    133
          2020-02-20    129
          ...
          2020-06-21      1
          2019-11-04      1
          2019-10-03      1
          2019-10-30      1
          2020-03-28      1
          Name: transaction_date, Length: 790, dtype: int64
```

```
In [229]: print(raw_data_nan['owner'].value_counts())
          print((raw_data_nan['owner'].astype(str) == 'nan').sum()
                + (raw_data_nan['owner'].astype(str) != 'nan').sum())
          (raw_data_nan['owner'].astype(str) == 'nan').sum()
```

```
joint      4418
self       2799
dependent   388
Name: owner, dtype: int64
14274
```

```
Out[229]: 6669
```

A significant amount of owner types are still nan, this could be looked at for missingness analysis

```
In [300]: ▶ print(raw_data_nan['ticker'].value_counts())
          (raw_data_nan['ticker'].astype(str) == 'nan').sum()
```

```
MSFT      239
AAPL      173
NTAP      124
TDDXX     122
FB        106
...
GME        1
PB         1
PSB        1
QRTEA      1
LSFYX      1
Name: ticker, Length: 2072, dtype: int64
```

Out[300]: 1147

there are around 1,000 nan's here in tickers... I'm thinking maybe they didn't invest particularly in stock but some sort of investment that doesn't have a ticker?

```
In [232]: ▶ print(raw_data_nan['asset_description'].value_counts())
          (raw_data_nan['asset_description'].astype(str) == 'nan').sum()
```

```
Microsoft Corporation      191
BLF FedFund                119
Apple Inc.                 113
Sunrun Inc.                93
Apple Inc                   61
...
Proctor & Gamble Company    1
Prudential Financial        1
Reata Pharmaceuticals, Inc - Class A stock  1
Republic Services           1
GrubHub Inc.                1
Name: asset_description, Length: 5000, dtype: int64
```

Out[232]: 4

only 4 missing here

```
In [233]: ▶ print(raw_data_nan['type'].value_counts())
          raw_data_nan['type'].value_counts().sum()
```

```
purchase      7428
sale_full     4563
sale_partial  2155
exchange      128
Name: type, dtype: int64
```

Out[233]: 14274

All types are present

```
In [234]: ▶ print(raw_data_nan['amount'].value_counts())
raw_data_nan['amount'].value_counts().sum()
```

```
$1,001 - $15,000          10054
$15,001 - $50,000         2258
$50,001 - $100,000        722
$100,001 - $250,000        545
$1,001 -                  242
$250,001 - $500,000        227
$500,001 - $1,000,000       142
$1,000,001 - $5,000,000     38
$1,000,000 +              28
$5,000,001 - $25,000,000    9
$1,000 - $15,000           4
$15,000 - $50,000          3
$50,000,000 +              1
$1,000,000 - $5,000,000     1
Name: amount, dtype: int64
```

Out[234]: 14274

Hmmm, not sure how to interpret \$1,001 - , but all amounts are present

I tried to webscrape because the site that hosts the data conveniently had all formatted names and party affiliation in one place but the get request never seemed to return the html I was looking for, bummer.

```
In [51]: raw_get = requests.get('https://housestockwatcher.com/summary_by_rep
raw_get.text
#stock_party_soup = bs4.BeautifulSoup(raw_get.text)
#stock_party_soup
```

```
Out[51]: '<!doctype html><html lang="en"><head><meta charset="utf-8"/><meta
name="viewport" content="width=device-width,initial-scale=1"/><meta
http-equiv="X-UA-Compatible" content="IE=edge,chrome=1"/><meta cont
ent="width=device-width,initial-scale=1,shrink-to-fit=no" name="vie
wport"/><meta name="title" content="House Stock Watcher - See What
Your Representative Is Trading"><meta name="description" content="U
pdated Daily - See the stock trades US Representatives are making a
s they are reported. Get notifications when new reports are uploade
d. Get insight now!"><meta property="og:type" content="website"><me
ta property="og:url" content="https://housestockwatcher.com"><meta
property="og:title" content="House Stock Watcher - See What Your Re
presentative Is Trading"><meta property="og:description" content="U
pdated Daily - See the stock trades US Representatives are making a
s they are reported. Get notifications when new reports are uploade
d. Get insight now!"><meta property="og:image" content="https://hou
sestockwatcher.com/promo1.png"><meta property="twitter:card" conten
t="summary_large_image"><meta property="twitter:url" content="http
s://housestockwatcher.com/"><meta property="twitter:title" content
="House Stock Watcher - See What Your Representative Is Trading"><m
eta property="twitter:description" content="Updated Daily - See the
stock trades US Representatives are making as they are reported. Ge
t notifications when new reports are uploaded. Get insight now!"><m
eta property="twitter:image" content="https://housestockwatcher.co
m/promo1.png"><title>House Stock Watcher - See What Your Representa
tive Is Trading</title><link rel="shortcut icon" href="/favicon.pn
g"/><link rel="apple-touch-icon" sizes="76x76" href="/apple-icon.pn
g"/><link rel="manifest" href="/manifest.json"/><link rel="styleshe
et" href="/tailwind.min.css"/><link href="https://unpkg.com/tailwin
dcss@^2/dist/tailwind.min.css" rel="stylesheet"><script async src
="https://www.googletagmanager.com/gtag/js?id=G-Z63X8GDP3C"></scrip
t><script>function gtag(){dataLayer.push(arguments)}window.dataLaye
r=window.dataLayer||[],gtag("js",new Date),gtag("config","G-Z63X8GD
P3C")</script><script async src="https://pagead2.googlesyndication.
com/pagead/js/adsbygoogle.js?client=ca-pub-2466331850146937" crosso
rigin="anonymous"></script><link href="/static/css/2.15bea899.chun
k.css" rel="stylesheet"><link href="/static/css/main.ce686fcf.chun
k.css" rel="stylesheet"></head><body class="text-blueGray-700 antia
liased"><noscript>You need to enable JavaScript to run this app.</n
oscript><div id="root"></div><script>!function(e){function t(t){for
(var n,f,l=t[0],a=t[1],c=t[2],p=0,s=[];p<l.length;p++)f=l[p],Objec
t.prototype.hasOwnProperty.call(o,f)&&o[f]&&s.push(o[f][0]),o[f]=0;
for(n in a)Object.prototype.hasOwnProperty.call(a,n)&&(e[n]=a[n]);f
or(i&&i(t);s.length;s.shift());return u.push.apply(u,c||[]),r()}f
unction r(){for(var e,t=0;t<u.length;t++){for(var r=u[t],n=!0,l=1;l
<r.length;l++){var a=r[l];0!==o[a]&&(n=!1)}n&&(u.splice(t--,1),e=f
(f.s=r[0]))}return e}var n={},o={1:0},u=[];function f(t){if(n[t])re
turn n[t].exports;var r=n[t]={i:t,l:!1,exports:{}};return e[t].call
(r.exports,r,r.exports,f),r.l=!0,r.exports}f.m=e,f.c=n,f.d=function
(e,t,r){f.o(e,t)||Object.defineProperty(e,t,{enumerable:!0,get:
r})},f.r=function(e){"undefined"!=typeof Symbol&&Symbol.toStringTag
&&Object.defineProperty(e,Symbol.toStringTag,{value:"Module"}),Obje
ct.defineProperty(e,"__esModule",{value:!0})},f.t=function(e,t){if
```



```
(1&t&&(e=f(e)),8&t)return e;if(4&t&&"object"===typeof e&&e&&e.__esModule)return e;var r=Object.create(null);if(f.r(r),Object.defineProperty(r,"default",{enumerable:!0,value:e}),2&t&&"string"!==typeof e)for(var n in e)f.d(r,n,function(t){return e[t]}.bind(null,n));return r},f.n=function(e){var t=e&&e.__esModule?function(){return e.default}:function(){return e};return f.d(t,"a",t),t},f.o=function(e,t){return Object.prototype.hasOwnProperty.call(e,t)},f.p="/";var l=this["webpackJsonphouse-stock-watcher-frontend"]=this["webpackJsonphouse-stock-watcher-frontend"]||[],a=l.push.bind(l);l.push=t,l=l.slice();for(var c=0;c<l.length;c++)t(l[c]);var i=a;r()}([])</script><script src="/static/js/2.d782ad26.chunk.js"></script><script src="/static/js/main.8db78f46.chunk.js"></script></body></html>'
```

Now that raw data is all cleaned up lets merge with our political party dataset which has congressman names and party affiliation. I was able to copy and paste the text from this website https://ballotpedia.org/List_of_current_members_of_the_U.S._Congress (https://ballotpedia.org/List_of_current_members_of_the_U.S._Congress).

right into google sheets and then generate a csv but I still had to clean up around 90 names manually. Some had middle initials, some former members were not included and there were some that had titles like Mr. Mrs. and Dr. This process took quite a while and actually made me appreciate/respect the data cleaning process more.

```

In [303]:  congress_party = pd.read_csv('congress_parties.csv')
           congress_party.shape

           congress_party = pd.read_csv('congress_parties.csv')
           #split name so we can extract last name
           congress_party['last_name'] = congress_party['full_name']
           .apply(str.split, ' ')
           #extract last name and set it to new column
           congress_party['last_name'] = congress_party['last_name']
           .apply(lambda x: x[-1])
           print(congress_party['last_name'].nunique())
           congress_party.nunique()

           #merge and then verify all names have been mapped to political party
           merged = raw_data_nan.merge(congress_party, left_on = 'representative
                                     right_on='full_name', how='left')
           merged.groupby('party').count().sum()#[merged['full_name'].astype(st

```

502

```

Out[303]: disclosure_year      14274
          disclosure_date      14274
          transaction_date     14274
          owner                7605
          ticker              13127
          asset_description    14270
          type                14274
          amount              14274
          representative       14274
          district            14274
          ptr_link             14274
          cap_gains_over_200_usd 14274
          full_name            14274
          last_name            14274
          dtype: int64

```

```

In [305]:  merged['ticker'].value_counts().sum()

```

```

Out[305]: 13127

```

About 1000 tickers missing, maybe

While it says there are 502 unique congress members part of the data set there is actually a bit less, when manually cleaning name mappings I noticed some transactions had different names for the same people, for example same person but one identity might have mr. or a middle initial while another just has first and last name. When looking at value counts we see they all add up to the original size of the data set except for owner which has a significant amount of values missign almost half and asset description missing 4. For our missingness question we will probably be inquiring further about the mechanism behind the owner data missing.

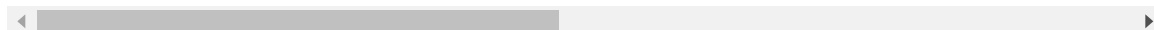
Now that we have our merged and cleaned dataset let's perform some univariate analysis and take a look into some of the individual columns

In [236]: raw_data_nan

Out[236]:

	disclosure_year	disclosure_date	transaction_date	owner	ticker	asset_description
0	2021	2021-10-04	2021-09-27	joint	BP	BP plc
1	2021	2021-10-04	2021-09-13	joint	XOM	Exxon Mobil Corporation
2	2021	2021-10-04	2021-09-10	joint	ILPT	Industrial Logistics Properties Trust - Common...
3	2021	2021-10-04	2021-09-28	joint	PM	Phillip Morris International Inc
4	2021	2021-10-04	2021-09-17	self	BLK	BlackRock Inc
...
14269	2020	2020-06-10	2020-04-09	NaN	SWK	Stanley Black & Decker, Inc.
14270	2020	2020-06-10	2020-04-09	NaN	USB	U.S. Bancorp
14271	2020	2020-06-10	2020-03-13	NaN	BMJ	Bristol-Myers Squibb Company
14272	2020	2020-06-10	2020-03-13	NaN	LLY	Eli Lilly and Company
14273	2020	2020-06-10	2020-03-13	NaN	DIS	Walt Disney Company

14274 rows × 12 columns



Let's take a quick look at the range of the year in which data was gathered

In [237]: raw_data_nan['disclosure_year'].value_counts()

Out[237]: 2020 7379
2021 5520
2022 1375
Name: disclosure_year, dtype: int64

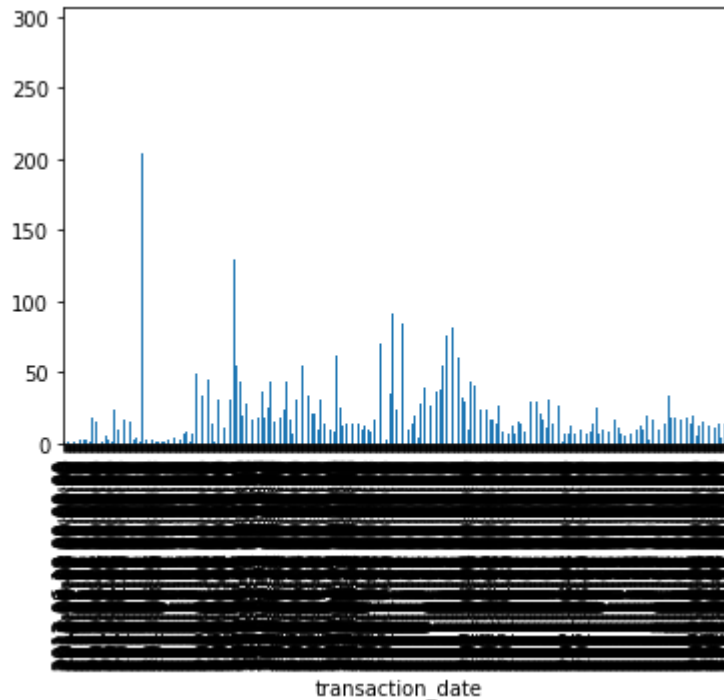
The years are quite relevant and given the turbulence of the markets over the last recent year this should be an interesting analysis

lets look at the volume of trades by date

```
In [251]: ▶ date_group = raw_data_nan.groupby('transaction_date').count()

date_group['disclosure_year'].plot(kind='bar')
#date_group.plot(kind='bar', x='transaction_year')
```

Out[251]: <AxesSubplot:xlabel='transaction_date'>



Well that doesn't look pretty at all, even with the messed up x axis we can kind of see some huge spikes this may be because the transaction dates were not properly recorded? It's weird that there is a massive spike early on and seemingly nothing around the massive spike. I should probably create a column that contains the month, but lets try the years instead

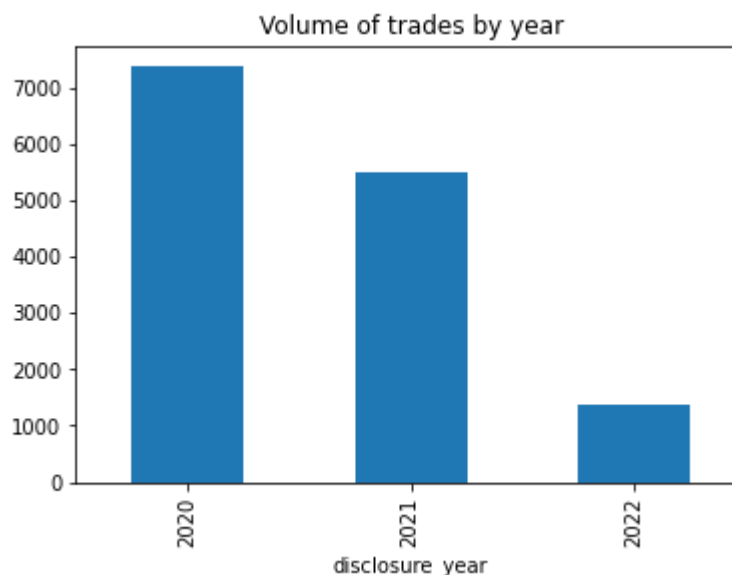
```
In [270]: ▶ date_group = raw_data_nan.groupby('transaction_date').count()  
          date_group['disclosure_date'].sort_values(ascending = False)[:20]
```

```
Out[270]: transaction_date  
2020-03-18    292  
2019-06-24    204  
2021-02-16    194  
2020-11-13    133  
2020-04-02    129  
2020-02-20    129  
2021-02-11    124  
2020-10-22     98  
2020-11-05     91  
2020-03-23     90  
2020-11-19     84  
2021-02-05     81  
2020-04-29     77  
2020-01-13     76  
2021-01-28     75  
2020-08-07     71  
2020-10-16     70  
2020-03-27     70  
2021-01-15     69  
2021-02-17     69  
Name: disclosure_date, dtype: int64
```

We can sort of get some sense of the months now

```
In [254]: ▶ date_group = raw_data_nan.groupby('disclosure_year').count()  
          date_group['disclosure_date'].plot(kind='bar',  
                                             title='Volume of trades by year')
```

```
Out[254]: <AxesSubplot:title={'center':'Volume of trades by year'}, xlabel='disclosure_year'>
```

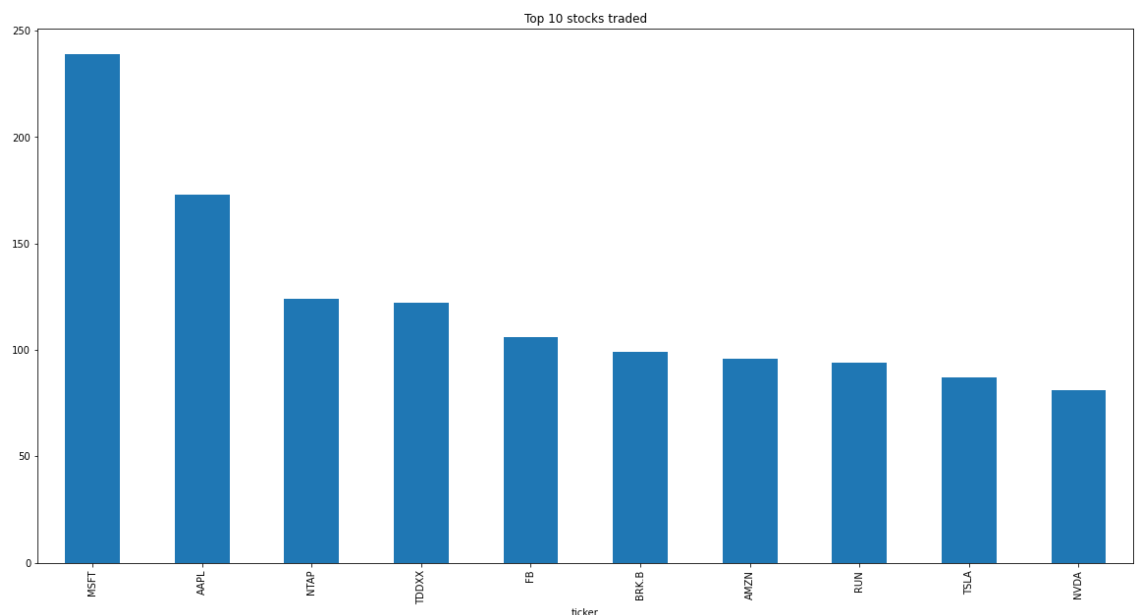


Okay this is pretty interesting we see that there were way more trades in 2020 and 2021 compared to 2022, maybe congressmembers knew something about the direction the economy was heading? 2022 has been a bloodbath so it makes sense why there is less trade compared to the booming markets of 2020 and 2021, later we'll look at whether they were selling or buying.

we can take a quick look at the hottest stocks congressmembers were trading, let's look at the top 10

```
In [291]: group_stocks = raw_data_nan.groupby('ticker').count()  
          .sort_values(by='disclosure_year', ascending=False)  
          plt.figure(figsize=(20, 10))  
          group_stocks['disclosure_year'][:10].plot(kind='bar',  
                                                    title='Top 10 stocks trad
```

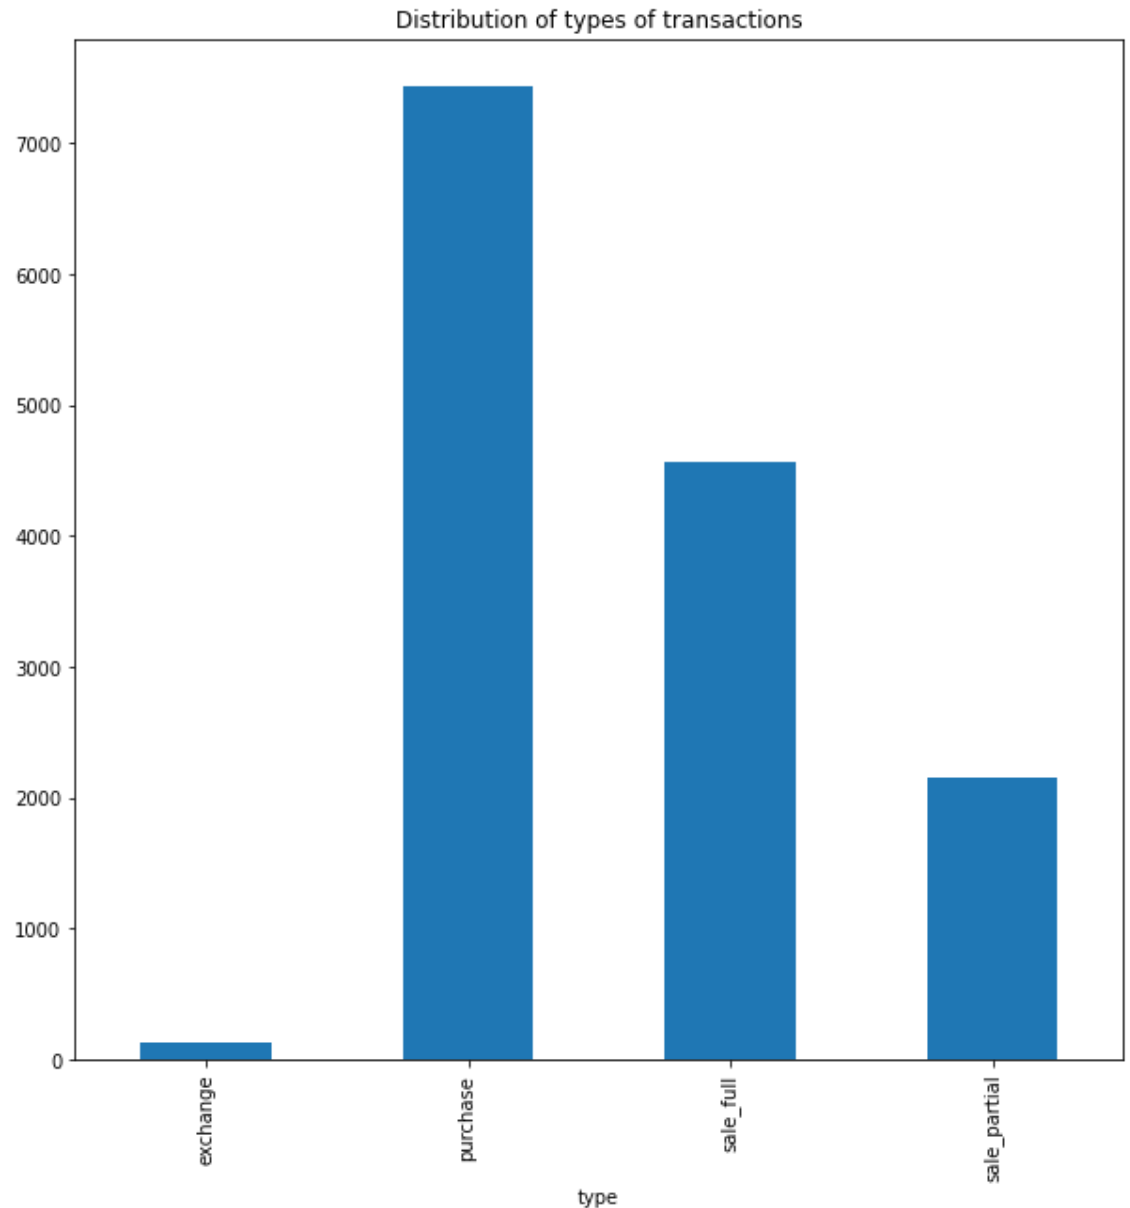
```
Out[291]: <AxesSubplot:title={'center': 'Top 10 stocks traded'}, xlabel='ticker'>
```



lets take a look at the distribution of types transactions

```
In [290]: group_types = raw_data_nan.groupby('type').count()  
plt.figure(figsize=(10, 10))  
group_types['disclosure_year'].plot(kind='bar',  
                                     title= 'Distribution of types of
```

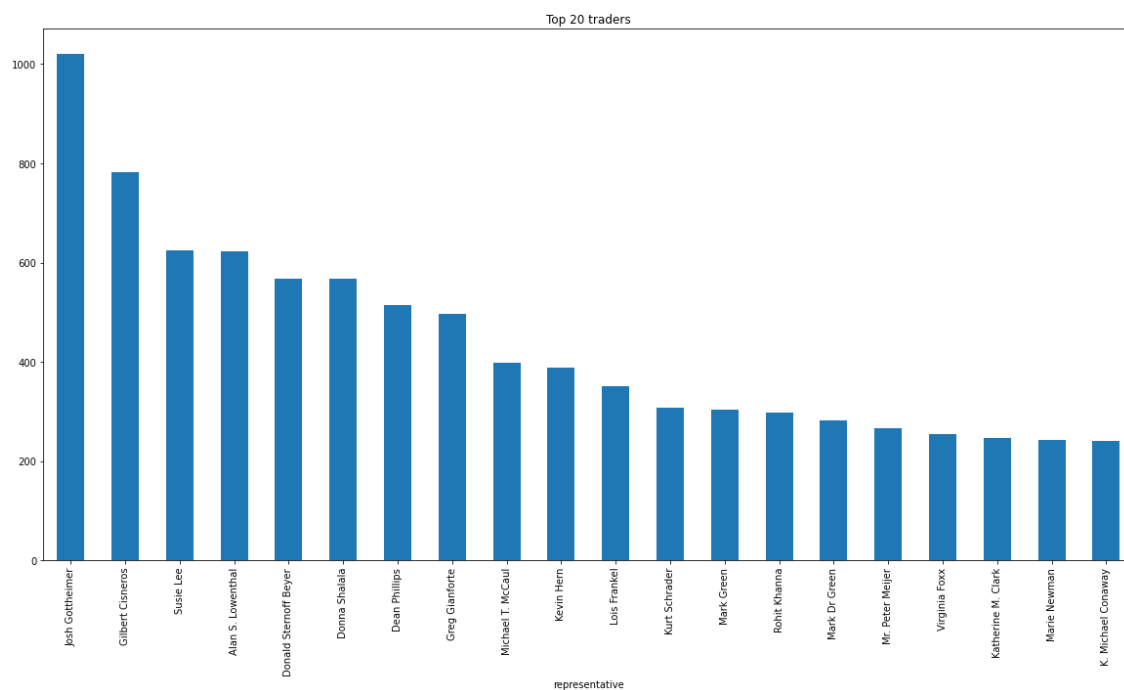
```
Out[290]: <AxesSubplot:title={'center': 'Distribution of types of transactions'}, xlabel='type'>
```



Let's check real quick to see who is making the most trades!

```
In [286]: group_rep = raw_data_nan.groupby('representative').count()
          .sort_values(by='disclosure_year',ascending=False)
          plt.figure(figsize=(20, 10))
          group_rep['disclosure_year'][:20].plot(kind='bar',
                                                title= 'Top 20 traders')
```

```
Out[286]: <AxesSubplot:title={'center':'Top 20 traders'}, xlabel='representative'>
```

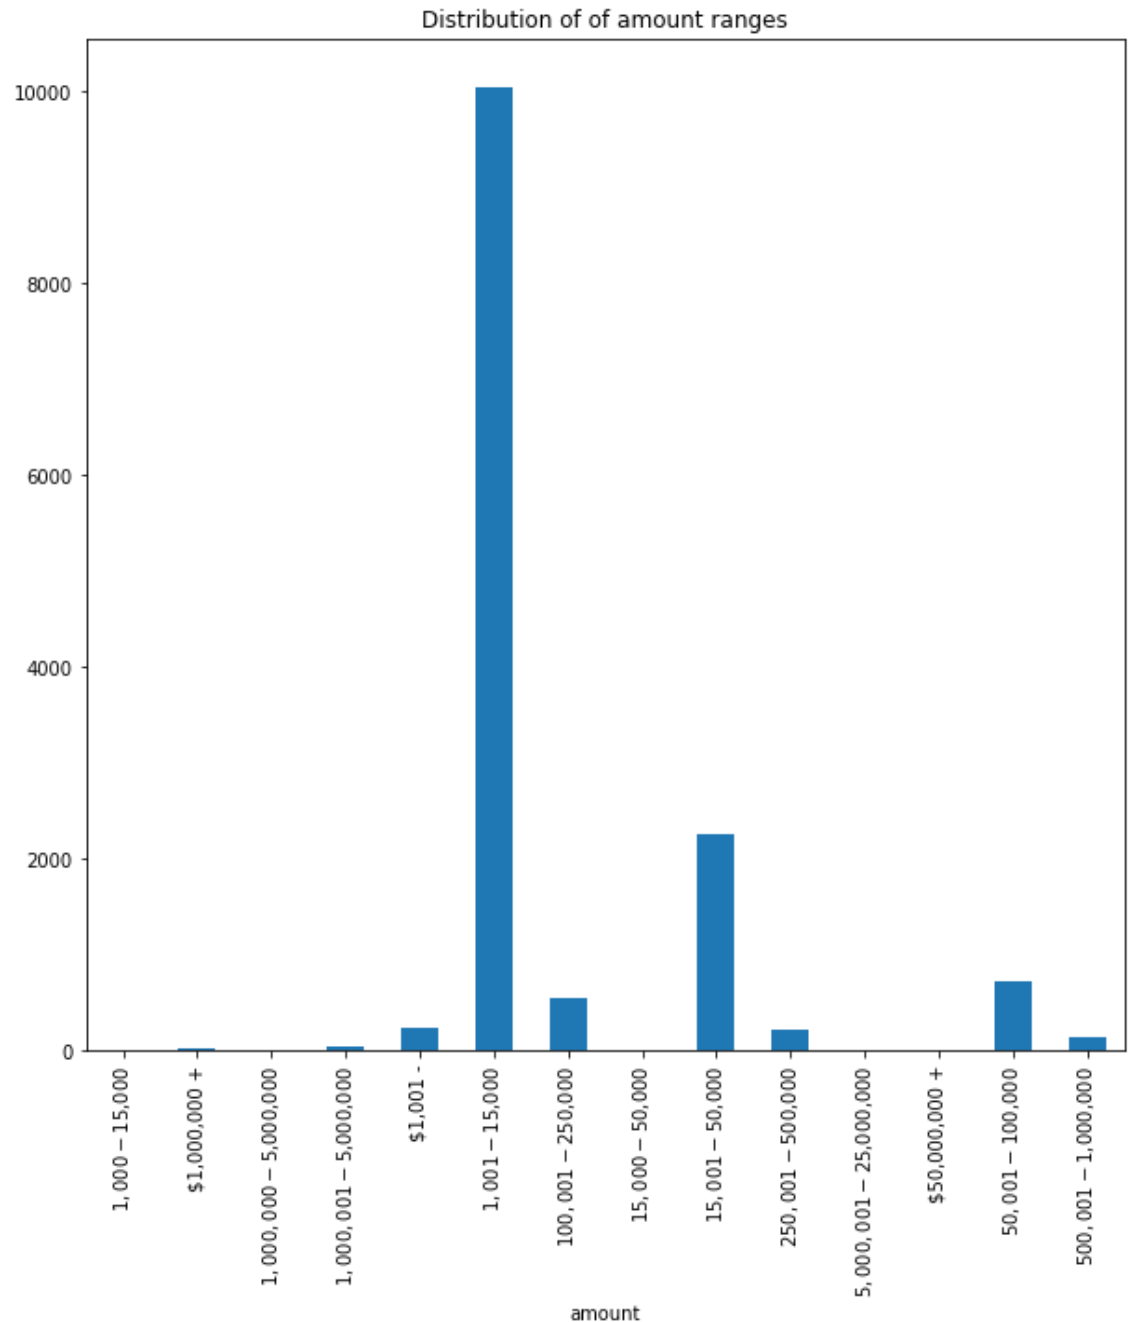


Okay we have some names but maybe in bivariate analysis we can explore how these names relate to party affiliation and look at proportion

before we dive into bivariate and aggregate analysis based on party affiliation let's look at distribution of amount ranges


```
In [283]: ▶ group_amount = raw_data_nan.groupby('amount').count()  
plt.figure(figsize=(10, 10))  
group_amount['disclosure_year'].plot(kind='bar',  
                                     title= 'Distribution of of amount
```

```
Out[283]: <AxesSubplot:title={'center': 'Distribution of of amount ranges'}, x  
label='amount'>
```



We see most trades were in \$1,000 to \$15,000 but there were a couple that were above \$50,000 let's check real quick how many.

```
In [280]: ▶ group_amount = raw_data_nan.groupby('amount').count()
          group_amount['owner']
```

```
Out[280]: amount
$1,000 - $15,000          1
$1,000,000 +             28
$1,000,000 - $5,000,000   0
$1,000,001 - $5,000,000   34
$1,001 -                 232
$1,001 - $15,000         5201
$100,001 - $250,000       272
$15,000 - $50,000         1
$15,001 - $50,000        1216
$250,001 - $500,000       136
$5,000,001 - $25,000,000   8
$50,000,000 +             1
$50,001 - $100,000        373
$500,001 - $1,000,000     102
Name: owner, dtype: int64
```

```
In [279]: ▶ 28 + 34 +272 + 136+8+1+373+102
```

```
Out[279]: 954
```

As a proportion of all trades these high value trades account for ...

```
In [281]: ▶ (28 + 34 +272 + 136+8+1+373+102) / group_amount['owner'].sum()
```

```
Out[281]: 0.12544378698224853
```

To explore bivariate distributions we'll mainly plot conditional plots based on part affiliation and pivot tables will be nice too let's use out merged data set

```
In [289]: merged.head()
```

Out[289]:

	disclosure_year	disclosure_date	transaction_date	owner	ticker	asset_description	
0	2021	10/04/2021	2021-09-27	joint	BP	BP plc	pl
1	2021	10/04/2021	2021-09-13	joint	XOM	Exxon Mobil Corporation	pl
2	2021	10/04/2021	2021-09-10	joint	ILPT	Industrial Logistics Properties Trust - Common...	pl
3	2021	10/04/2021	2021-09-28	joint	PM	Phillip Morris International Inc	pl
4	2021	10/04/2021	2021-09-17	self	BLK	BlackRock Inc	sale

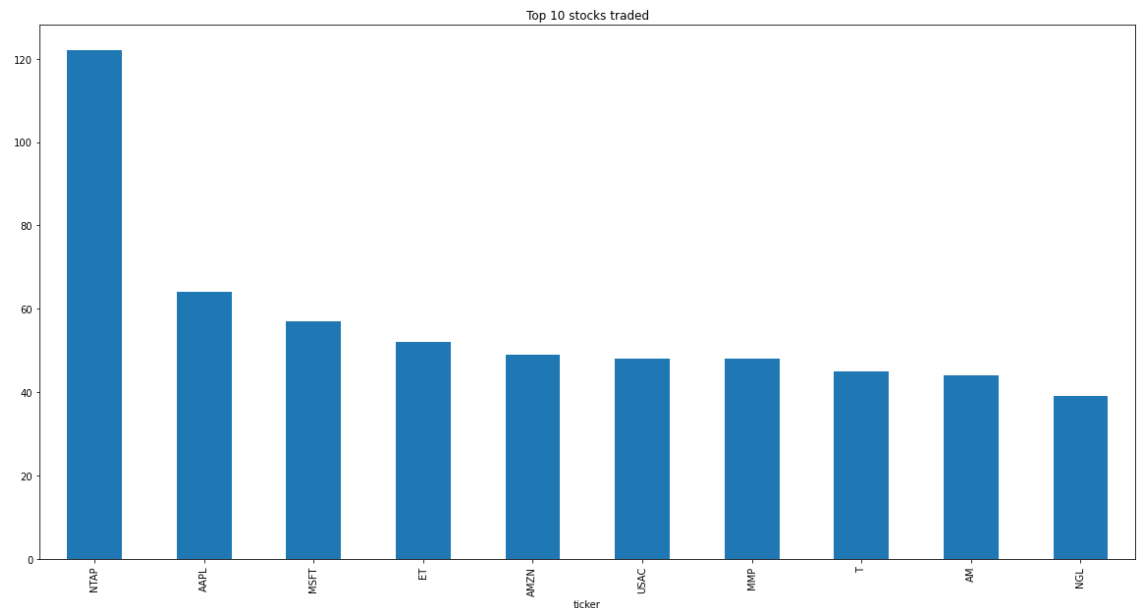


```
lets revisit the top 10 stocks based on party affiliation
```

```
In [306]: group_stocks = merged[merged['party'] == 'Republican']
          .groupby('ticker').count().sort_values(by='disclosure_year',
                                                  ascending=False)

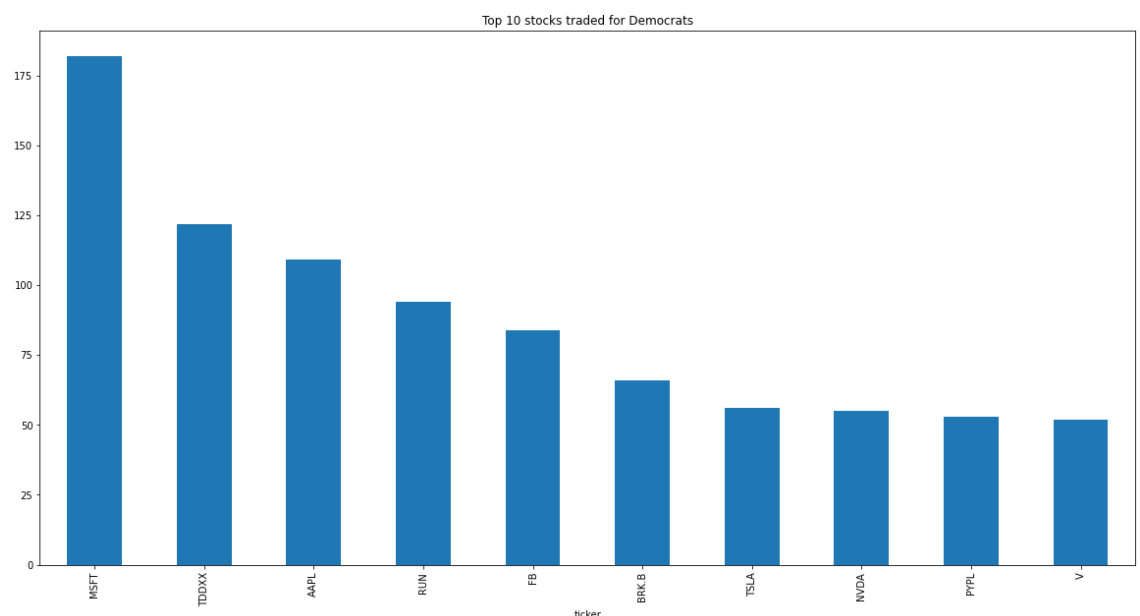
plt.figure(figsize=(20, 10))
group_stocks['disclosure_year'][:10].plot(kind='bar',
                                          title='Top 10 stocks trad
```

```
Out[306]: <AxesSubplot:title={'center':'Top 10 stocks traded'}, xlabel='ticke
r'>
```



```
In [308]: group_stocks = merged[merged['party'] == 'Democratic']
          .groupby('ticker').count().sort_values(by='disclosure_year', asce
plt.figure(figsize=(20, 10))
group_stocks['disclosure_year'][:10].plot(kind='bar',
                                          title='Top 10 stocks trad
```

```
Out[308]: <AxesSubplot:title={'center':'Top 10 stocks traded for Democrats'},
xlabel='ticker'>
```



This is pretty interesting among the top 10 only apple and microsoft are shared amongst the top 10 stocks, We see for republicans there are sort of older well established tech and traditional energy while democrats are trading new tech, solar energy, and blackrock + berkshire

Let's see which party makes more trades

```
In [ ]: merged.sort_values(by='pivot_table(index='')
```

```
In [331]: group_rep = raw_data_nan.groupby('representative').count()
          .sort_values(by='disclosure_year', ascending=False)
          all_trades = pd.DataFrame(group_rep['disclosure_year'])
          .merge(congress_party, left_index = True, right_on='full_name', h
          all_trades.groupby('party')['disclosure_year'].sum() / raw_data_nan.
```

```
Out[331]: party
Democratic    0.638434
Republican    0.361566
Name: disclosure_year, dtype: float64
```

```
In [ ]: merged.pivot_table(index='party', columns='')
```

Okay so we observe from the data that Democrats tend to trade more, how extreme is this value, could this be due to chance? I think I'll hypothesis test if the distribution of trades comes from a null distribution of 50% / 50% from each party, we will visit this later

Lets look at the type of trades congressmembers make based on party affiliation

```
In [333]: types = merged.pivot_table(index='party', columns = 'type',
          values='representative', aggfunc='count')
          types
```

```
Out[333]:
```

type	exchange	purchase	sale_full	sale_partial
party				
Democratic	57	4373	2925	1758
Republican	71	3055	1638	397

We can clearly see much more buying and much more selling from democrats but lets dig a bit deeper

Okay this is interesting, let's extend this to also see the types of trades within years as proportions

```
In [385]:  types = merged.pivot_table(index='party',columns = ['type','disclosure_year'],values='representative',aggfunc='count')
types / types.sum(axis=0)
```

Out[385]:

type	exchange			purchase			sale_full	
disclosure_year	2020	2021	2022	2020	2021	2022	2020	2021
party								
Democratic	0.43956	0.518519	0.3	0.601844	0.583416	0.543417	0.713609	0.548747
Republican	0.56044	0.481481	0.7	0.398156	0.416584	0.456583	0.286391	0.451253

We can see that the majority of purchases were conducted in 2020 and 2021 amongst both parties and that democratic congress members participated in much more selling activity both full and partialy. Okay we can see they participate in more raw number of trades but what about the actual amount of money?

```
In [388]: types = merged.pivot_table(columns='party', index = ['type', 'amount'],
                                     values='representative', aggfunc='count')
types
```

Out[388]:

	party	Democratic	Republican
type	amount		
exchange	1, 000, 001–5,000,000	NaN	1.0
	\$1,001 -	1.0	NaN
	1, 001–15,000	44.0	17.0
	100, 001–250,000	1.0	9.0
	15, 001–50,000	6.0	15.0
	250, 001–500,000	2.0	7.0
	50, 001–100,000	3.0	21.0
	500, 001–1,000,000	NaN	1.0
purchase	\$1,000,000 +	9.0	NaN
	1, 000, 001–5,000,000	17.0	2.0
	\$1,001 -	88.0	5.0
	1, 001–15,000	3567.0	1626.0
	100, 001–250,000	47.0	230.0
	15, 001–50,000	461.0	814.0
	250, 001–500,000	36.0	59.0
	5, 000, 001–25,000,000	4.0	NaN
sale_full	50, 001–100,000	90.0	301.0
	500, 001–1,000,000	54.0	18.0
	1, 000–15,000	NaN	3.0
	\$1,000,000 +	12.0	1.0
	1, 000, 000–5,000,000	NaN	1.0
	1, 000, 001–5,000,000	11.0	4.0
	\$1,001 -	21.0	1.0
	1, 001–15,000	2267.0	826.0
	100, 001–250,000	62.0	156.0
	15, 000–50,000	NaN	2.0
	15, 001–50,000	369.0	415.0
	250, 001–500,000	52.0	58.0
	5, 000, 001–25,000,000	2.0	1.0
	\$50,000,000 +	NaN	1.0
	50, 001–100,000	102.0	144.0
	500, 001–1,000,000	27.0	25.0

party	type	amount	Democratic	Republican
sale_partial	1, 000–15,000		NaN	1.0
	\$1,000,000 +		6.0	NaN
	1, 000, 001–5,000,000		3.0	NaN
	\$1,001 -		126.0	NaN
	1, 001–15,000		1472.0	235.0
	100, 001–250,000		20.0	20.0
	15, 000–50,000		NaN	1.0
	15, 001–50,000		80.0	98.0
	250, 001–500,000		10.0	3.0
	5, 000, 001–25,000,000		2.0	NaN
	50, 001–100,000		23.0	38.0
	500, 001–1,000,000		16.0	1.0

There are some baffling findings, when looking at the \$500,000+ trades we clearly see that democrats engage in far more high value trades compared to republicans aside from the one republican trade for more than \$50,000,000.

Due to the coarse granularity of the amount variable it is definitely a limiting factor in unlocking this particular data set but it is still better than nothing.

Assessment of Missingness

During the EDA phase I noticed that the owner columns had almost half of it's values missing. Thinking about it, it could potentially be NMAR maybe if the transaction would appear shady based on a certian value it could look bad for the congressman, so they just don't report it at all. But we should conduct an appropriate MAR test to see if there is actually dependence on other columns, it could very well be the case the maybe party affiliation, type of trade, or amount of trade could be valid variables that influence whether the owner variable is missing so we'll conduct some MAR tests. Given the categorical nature of most columns, we will mostly use tvd as the test statistic. We will use 0.01 significance level.


```

In [372]: merged['owner_isnull'] = merged['owner'].isna()

emp_distributions = (
    merged
    .pivot_table(columns='owner_isnull', index='party',
                  aggfunc='size')
    .fillna(0)
    .apply(lambda x:x / x.sum())
)

emp_distributions.plot(kind='barh', title='distribution of party aff

observed_tvd = np.sum(np.abs(emp_distributions.diff(axis=1)
                             .iloc[:, -1])) / 2
print('observed tvd: ', observed_tvd)

n_repetitions = 500

merged_type = merged.copy()[['party', 'owner_isnull']]
tvds = []
for _ in range(n_repetitions):

    shuffled_types = (
        merged_type['party']
        .sample(replace=False, frac=1)
        .reset_index(drop=True)
    )

    shuffled = (
        merged_type
        .assign(**{'Shuffled Types': shuffled_types})
    )

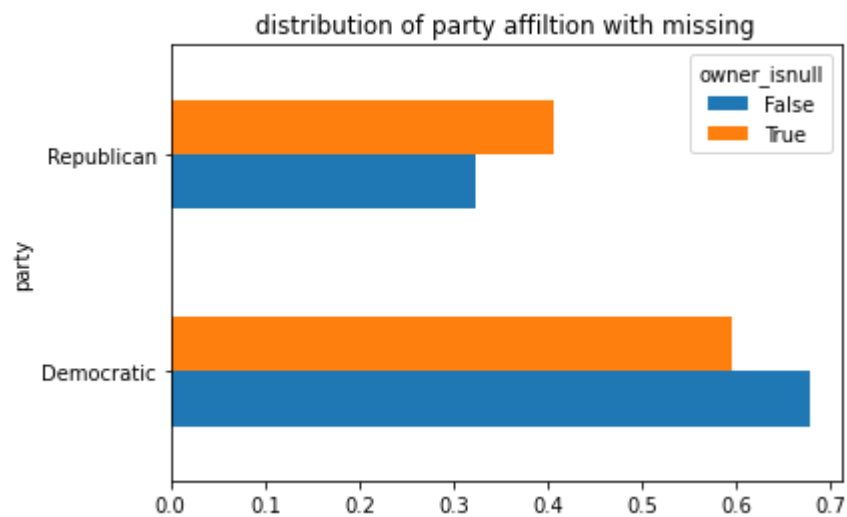
    # compute the tvd
    shuffled_emp_distributions = (
        shuffled
        .pivot_table(columns='owner_isnull', index='Shuffled Types',
                      values=None, aggfunc='size')
        .fillna(0)
        .apply(lambda x:x/x.sum())
    )

    tvd = np.sum(np.abs(shuffled_emp_distributions.diff(axis=1).iloc[
# add it to the list of results

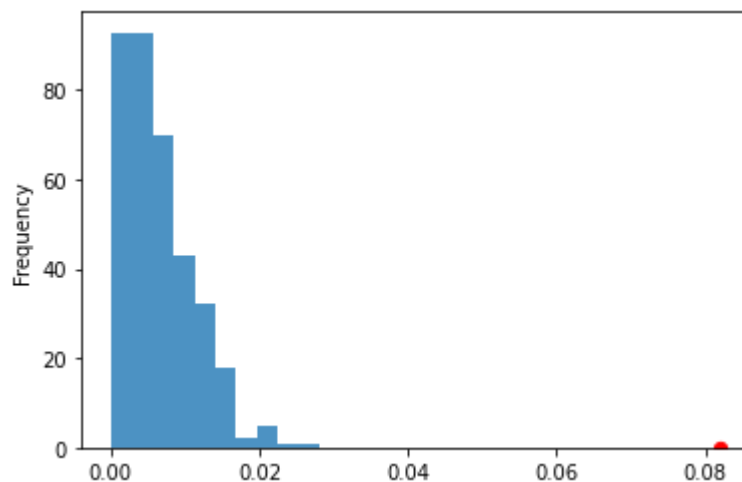
    tvds.append(tvd)

```

observed tvd: 0.08209972663413961



```
In [373]: ▶ pd.Series(tvds).plot(kind='hist', density=True, alpha=0.8)
plt.scatter(observed_tvd, 0, color='red', s=40);
```



```
In [374]: ▶ pval = np.mean(tvds >= observed_tvd)
pval
```

Out[374]: 0.0

```

In [375]: merged['owner_isnull'] = merged['owner'].isna()

emp_distributions = (
    merged
    .pivot_table(columns='owner_isnull', index='amount',
                  aggfunc='size')
    .fillna(0)
    .apply(lambda x:x / x.sum())
)

emp_distributions.plot(kind='barh',
                       title='distribution of amount with missingnes

observed_tvd = np.sum(np.abs(emp_distributions.diff(axis=1).iloc[:,-
print('observed tvd: ', observed_tvd)

n_repetitions = 500

payments_type = merged.copy()[['amount', 'owner_isnull']]
tvds = []
for _ in range(n_repetitions):

    shuffled_types = (
        payments_type['amount']
        .sample(replace=False, frac=1)
        .reset_index(drop=True)
    )

    shuffled = (
        payments_type
        .assign(**{'Shuffled Types': shuffled_types})
    )

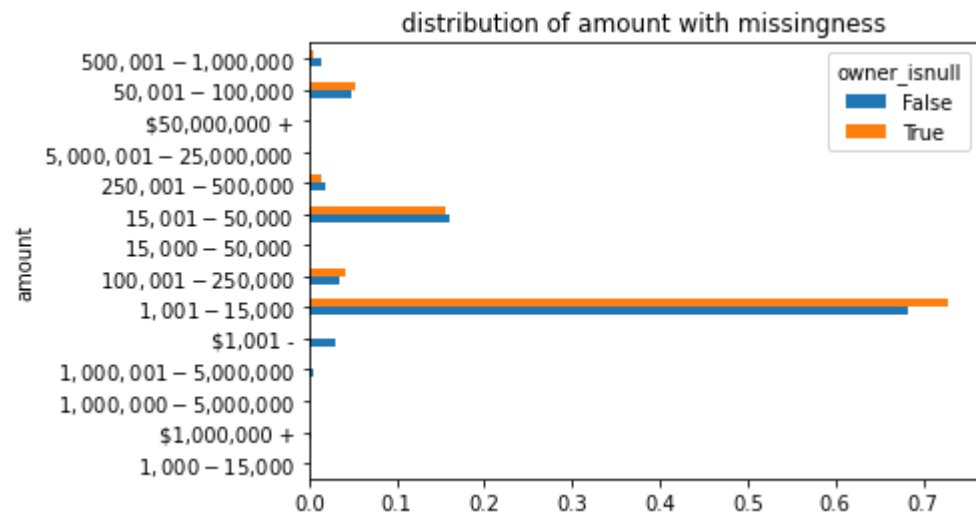
    # compute the tvd
    shuffled_emp_distributions = (
        shuffled
        .pivot_table(columns='owner_isnull', index='Shuffled Types',
                      values=None, aggfunc='size')
        .fillna(0)
        .apply(lambda x:x/x.sum())
    )

    tvd = np.sum(np.abs(shuffled_emp_distributions.diff(axis=1).iloc[
# add it to the list of results

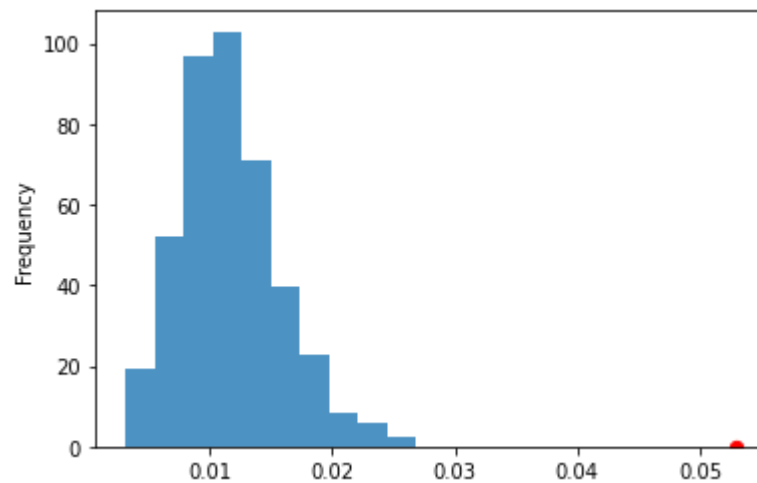
    tvds.append(tvd)

```

observed tvd: 0.05289456382573785



```
In [376]: ▶ pd.Series(tvds).plot(kind='hist', density=True, alpha=0.8)
plt.scatter(observed_tvd, 0, color='red', s=40);
```



```
In [377]: ▶ pval = np.mean(tvds >= observed_tvd)
pval
```

Out[377]: 0.0

```

In [381]: merged['owner_isnull'] = merged['owner'].isna()

emp_distributions = (
    merged
    .pivot_table(columns='owner_isnull', index='type', aggfunc='size'
    .fillna(0)
    .apply(lambda x:x / x.sum())
)

emp_distributions.plot(kind='barh',
                        title='distribution of type with missingness'

observed_tvd = np.sum(np.abs(emp_distributions.diff(axis=1)
                             .iloc[:, -1])) / 2
print('observed_tvd', observed_tvd)

n_repetitions = 500

merged_type = merged.copy()[['type', 'owner_isnull']]
tvds = []
for _ in range(n_repetitions):

    shuffled_types = (
        merged_type['type']
        .sample(replace=False, frac=1)
        .reset_index(drop=True)
    )

    shuffled = (
        merged_type
        .assign(**{'Shuffled Types': shuffled_types})
    )

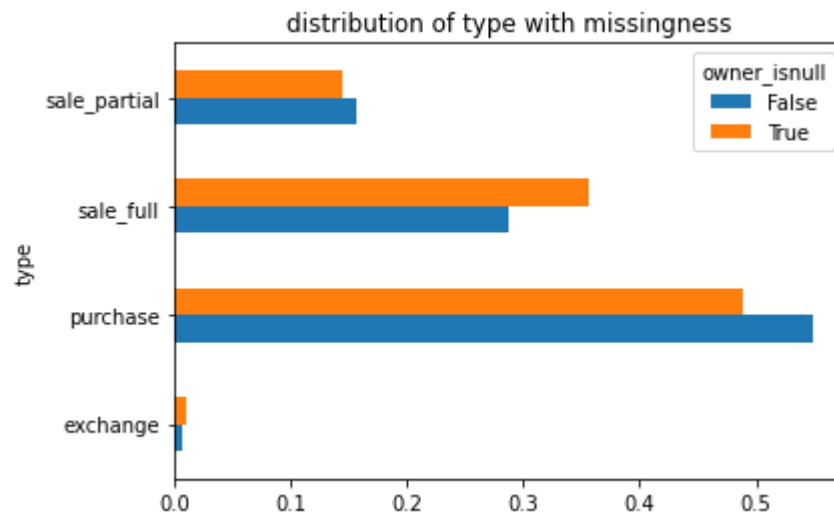
    # compute the tvd
    shuffled_emp_distributions = (
        shuffled
        .pivot_table(columns='owner_isnull', index='Shuffled Types',
                     values=None, aggfunc='size')
        .fillna(0)
        .apply(lambda x:x/x.sum())
    )

    tvd = np.sum(np.abs(shuffled_emp_distributions.diff(axis=1).iloc[
# add it to the list of results

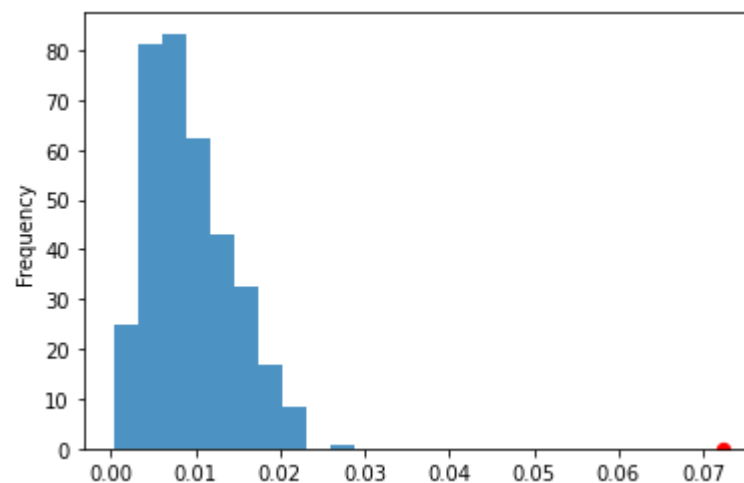
    tvds.append(tvd)

```

observed_tvd 0.07241542383242788



```
In [382]: ▶ pd.Series(tvds).plot(kind='hist', density=True, alpha=0.8)
           ▶ plt.scatter(observed_tvd, 0, color='red', s=40);
```



```
In [383]: ▶ pval = np.mean(tvds >= observed_tvd)
           ▶ pval
```

Out[383]: 0.0

Even with significance level of 0.01 we see from all three MAR permutation tests that the observed total variation distance among variables party, amount, and type are all much greater than the respective null distribution generated by permutation tests. In all three cases the null hypothesis that these tvd's come from the same distribution is rejected in favor of alternative hypothesis that they come from different distributions. Because these tvd's do not come from the same distributions that means there is a MAR dependence of owner on these variables. After obtaining MAR dependence verification from permutation testing we can think how intuitively these variables might influence the probability that the owner column is missing. In

terms of party it could be due to that fact that one party or the other might tend to report less or more often, for amount maybe the amount a very large transaction might have a higher probability to report owner because it could appear shady or might not want to disclose the ownership of the position based on size. For the type it could be maybe that buys or sells are over or under reported so as not to appear to buy or sell in a motivated fashion in which the ownership could hint at the scope of those involved.

Hypothesis Testing

Let's further explore the distribution behind the proportion of trades made between party affiliations, the test statistic of interest is the signed difference between the proportions, we'll use democrats - rep, my proposed null hypothesis is a world in which I would think both parties of interest would have an equal interest in trading so the proportion of trades from democrats is around 50% and the proportion of trades from republicans is around 50% we might expect our test statistic to be close to zero under our null. My alternative hypothesis is that we live in a world where democrats are more interested in trading, so the proportion of trades from democrats is above 50% and republicans trade less than 50% of the time yielding a test statistic that is positive because, $\text{dem \%} - \text{rep \%}$ would expect higher number of dem % trades. We will use 0.01 significance level.

```
In [339]: ▶ observed = all_trades.groupby('party')['disclosure_year'].sum() / ra
          observed = observed[0] - observed[1]
          observed
```

Out[339]: 0.2768670309653916

I think this test might be too easy but I have other homework and I think it answers an honest question so we'll go ahead

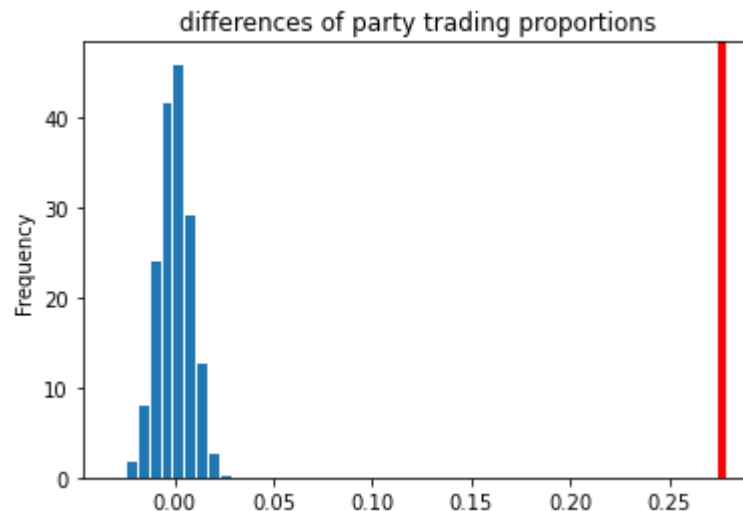
```
In [345]: ▶ N = 10000

          # 10000 times, we want to 'flip a coin' the size of original data se
          results = []
          for _ in range(N):
              simulation = np.random.choice([0, 1], p=[0.5, 0.5], size=merged.
              party_prop_diff = (simulation == 0).mean() - (simulation == 1).m
              results.append(party_prop_diff)
```

In [346]: `results`

```
Out[346]: [-0.013310914950259267,
-0.001961608518985525,
-0.005884825556956741,
0.0,
-0.004763906403250695,
-0.012189995796553221,
0.015973097940311043,
0.003502872355331421,
0.0025220680958386033,
-0.003502872355331421,
-0.005464480874316946,
-0.008547008547008517,
0.01429171920975203,
0.015973097940311043,
0.0044836766148241836,
-0.008126663864368833,
-0.008547008547008517,
0.0009808042594927624,
0.0030825276726915707,
0.000207607012200051]
```

In [347]: `pd.Series(results).plot(kind='hist',
density=True,
ec='w',
title='differences of party trading proportions',
plt.axvline(x=observed, color='red', linewidth=4);`



In [348]: `p_val = (results >= observed).mean()
p_val`

Out[348]: 0.0

It was pretty obvious from the graph but this difference in proportions is definitely not a feasible occurrence under the null hypothesis even with an extreme significance level of 0.01. Based on the p value of zero we will reject the null hypothesis that the interest of trading is evenly divided among republicans and democrats in favor of the alternative that democrats are more interested in trading than republicans.

